

ECE 661: Homework 4

Monday, 1 October 2018

Dr. Avinash Kak

Naveen Madapana

1. Interest Point Extraction

1.1 Definition

A point in the image is considered as an interest point if its pixel intensity is a local maxima or a local minima w.r.t x axis, y axis and the z axis (scale). An ideal interest point is translation, rotation and scale invariant. There are several interest point detectors, namely, 1. Scale Invariant Feature Transform (SIFT), 2. Sped Up Robust Features (SURF), 3. Harris corners, etc.

1.2 SIFT

The main steps in SIFT are summarized below.

1. The theorem that directly relates Laplacian of Gaussian (LoG) with Difference of Gaussian (DoG) is central to SIFT and SURF. It is given as follows. It means that LoG can be approximated by estimating the DoG pyramid (DoG at various scales)

$$\frac{\partial f(x, y, \sigma)}{\partial \sigma} = \sigma \Delta^2 f(x, y, \sigma) \quad (1)$$

2. DoG pyramid consists of scale sapce representation of an image at several octaves, i.e. $\sigma_0, 2\sigma_0, 4\sigma_0, 8\sigma_0, \dots$. Further, each level, multiple scales at minor variations of standard deviation are considered i.e. for instance at σ_0 , the variations are $\sigma_0 + d\sigma_0, \sigma_0 + 2d\sigma_0, \sigma_0 + 3d\sigma_0, \sigma_0 + 4d\sigma_0, \dots$. All of this constitutes the DoG pyramid.
3. Given the DoG pyramid, we mark interest point as the pixels that are either a local maxima or a local minima i.e. the intensity of an interest point is larger than its 27 neighbors (9 - same scale, 9 - a scale above, 9 - a scale below).
4. As σ increases, the spatial resolution gets coarser. As a result, the image needs to be downsampled by half at each level of space representation. In order to find the corresponding points of local maxima/minima in the original image, following Taylor's expansion is used:

$$D(\bar{X}) \approx D(\bar{X}_0) + J^T(\bar{X}_0)\bar{X} + \frac{1}{2}\bar{X}^T H(X_0)\bar{X} \quad (2)$$

Where, J is the Jacobian matrix (gradient vector) evaluated at the points of extrema. H is the Hessian matrix (of double derivatives) evaluated at \bar{X}_0 . Now, we will set $\frac{\partial D(\bar{X})}{\partial \bar{X}} = 0$. This gives the optimal $\bar{X} = -H^{-1}(\bar{X}_0)J(\bar{X}_0)$

5. Due to the errors in numerical precision and sampling, the value of $D(\bar{X}) < 0.1$ is thresholded to find more local extrema. We further make sure that the interest points are as far away from each other as possible.
6. Next, we find the dominant orientation at each interest point and consider a rectangle of size 8 x 16 in the dominant orientation at the interest point, and find a 128 dimensional descriptor vector at each interest point.

1.3 HARRIS Corners

In HARRIS corner detection, we aim to identify the points in the image at which the intensity change drastically both in x and y directions are considered as interest points. It utilizes the gradients in x and y computed using either haar filters or a simple Sobel operator. Further, this technique is applied at a particular scale. In contrast to SIFT, this technique is not invariant to scale. Nevertheless, these interest points are translation and rotation invariant.

Consider the following Haar filters to compute the gradients. The kernels are convolved across the images to obtain the gradients in x and y directions at a particular scale.

$$\frac{\partial}{\partial x} = \begin{bmatrix} -1 & -1 & -1 & 1 & 1 & 1 \\ -1 & -1 & -1 & 1 & 1 & 1 \\ -1 & -1 & -1 & 1 & 1 & 1 \\ -1 & -1 & -1 & 1 & 1 & 1 \\ -1 & -1 & -1 & 1 & 1 & 1 \\ -1 & -1 & -1 & 1 & 1 & 1 \end{bmatrix}$$

$$\frac{\partial}{\partial y} = \begin{bmatrix} -1 & -1 & -1 & -1 & -1 & -1 \\ -1 & -1 & -1 & -1 & -1 & 1 \\ -1 & -1 & -1 & -1 & -1 & -1 \\ 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

Next, a $5\sigma \times 5\sigma$ sized region is considered. Matrix C is estimated using this region at every point in the image. C is given below. Note that Σ is the summation over all of the pixels in the region.

$$C = \begin{bmatrix} \Sigma(d_x)^2 & \Sigma dxdy \\ \Sigma dxdy & \Sigma(dy)^2 \end{bmatrix}$$

Now, the points where the matrix C is non singular are considered as interest points. When a matrix is non singular, the eigen values will be non zero. Further we note that,

$$\text{trace}(C) = \lambda_1 + \lambda_2 = \Sigma(d_x)^2 + \Sigma(d_y)^2 \quad (3)$$

$$\det(C) = \lambda_1 \lambda_2 = \Sigma(d_x)^2 \Sigma(d_y)^2 - \Sigma(d_x d y)^2 \quad (4)$$

Let $r = \frac{\lambda_2}{\lambda_1}$, where λ_1 and λ_2 are eigenvalues and $\lambda_1 \geq \lambda_2$. Now,

$$\beta = \frac{\det(C)}{\text{trace}(C)^2} = \frac{\lambda_1 \lambda_2}{(\lambda_1 + \lambda_2)^2} = \frac{r}{(1+r)^2} \quad (5)$$

For Harris corner detection, we threshold $\beta > 0.08$, to find the interest points. Next, we perform non maxima suppression to reduce the false positives.

At each interest point, a matrix of size 31×31 is considered which is centered at the interest point. A descriptor vector is considered as an array of pixel intensities in the 31×31 region.

Establishing Point Correspondences

Let u and v be the vectors of same length (k). Let μ_u and μ_v be the mean of u and v respectively.

2.1 Sum Squared Differences (SSD)

$$SSD(u, v) = \sum_i (u_i - v_i)^2 \quad (6)$$

2.2 Normalized Cross Correlation (NCC)

$$NCC(u, v) = \frac{\sum_i (u_i - \mu_u)(v_i - \mu_v)}{\sqrt{\sum_i (u_i - \mu_u)^2} \sqrt{\sum_i (v_i - \mu_v)^2}} \quad (7)$$

The results obtained by using NCC are better than SSD as the former is mean normalized and is not sensitive to the uniform increase/decrease in the intensities. In experiments, the value of threshold NCC ranges anywhere from 0.97 to 0.999.

Experiments

SIFT (opencv), SURF (opencv) and HARRIS corner detection (own code) were used to find correspondences between four pairs of images. The first two pairs are given in the class, and the other two images were taken by me.

Further, the Harris corner detection is done two different scales (1.414 and 2.00). It was found that the quality of point correspondences is higher for higher scales due to smoothing. Also, SSD performs better than NCC unless the average intensity and scale of the image does not change.

SURF provided comparatively large number of interest points at the same threshold of NCC. Hence higher NCC thresholds are used to reduce the number of interest points.

Pair 1

Original Images



(a) Image 1

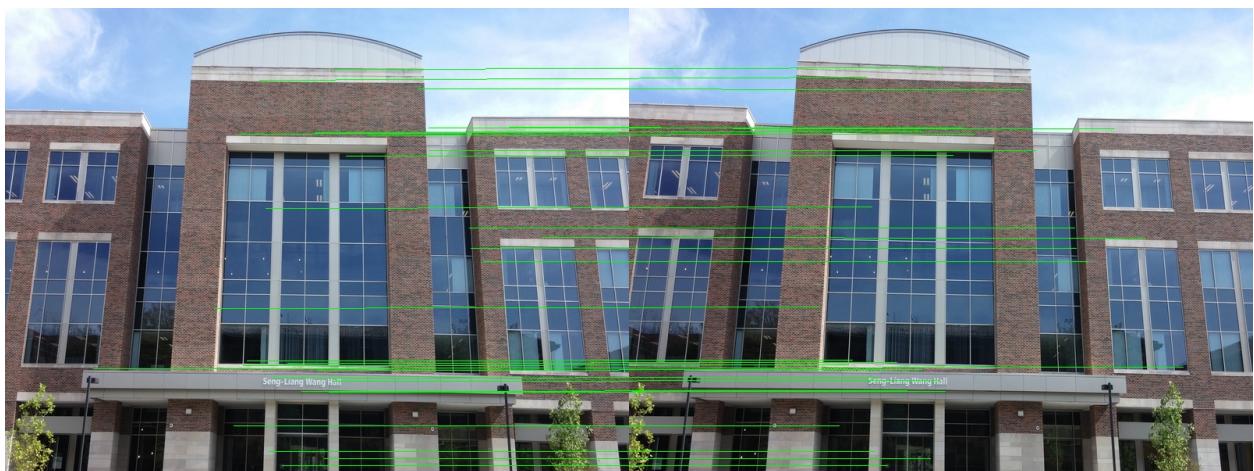


(b) Image 2

SIFT (thresh = 0.990)



SIFT (thresh = 0.99)



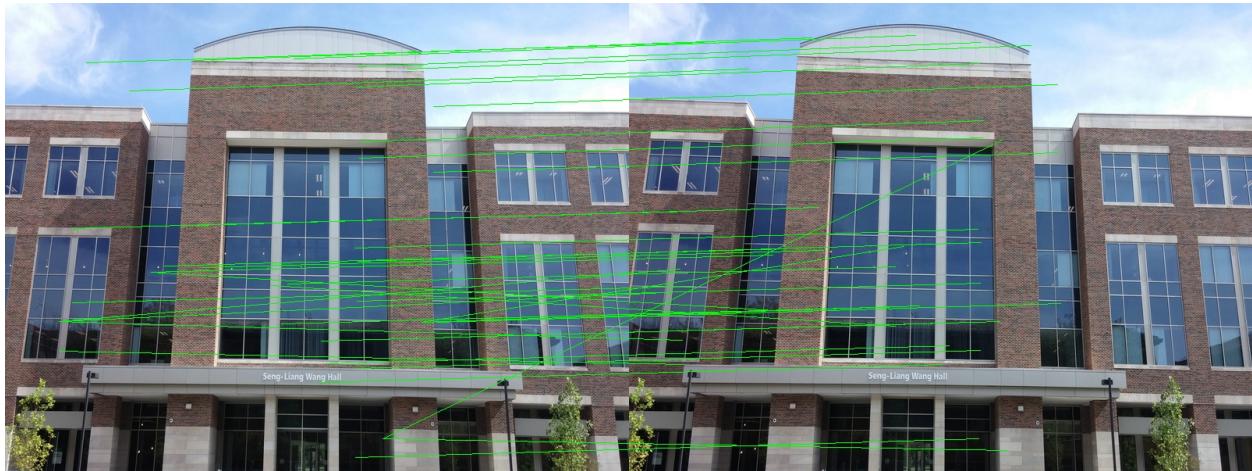
SURF (thresh = 0.9994)



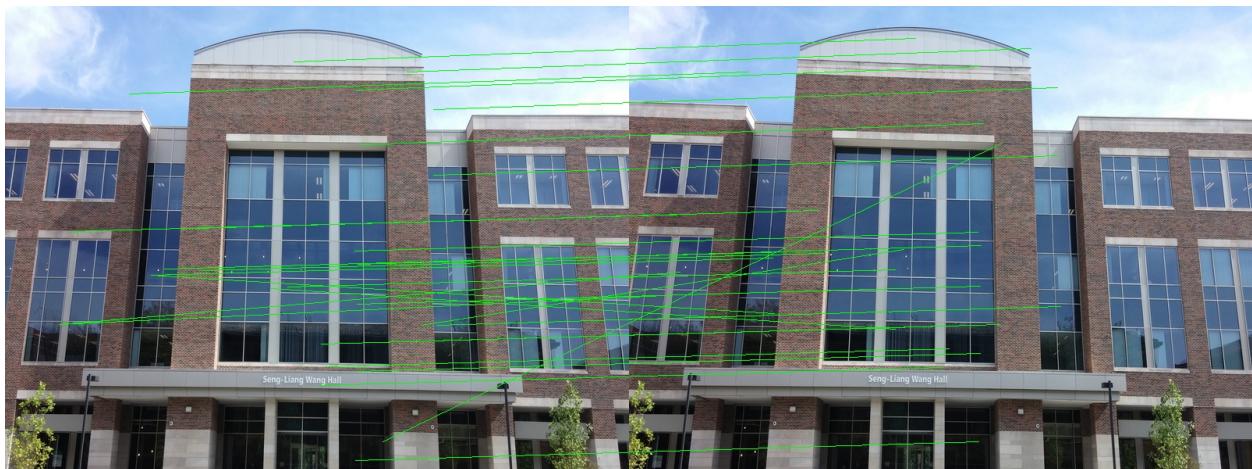
SURF (thresh = 0.9998)



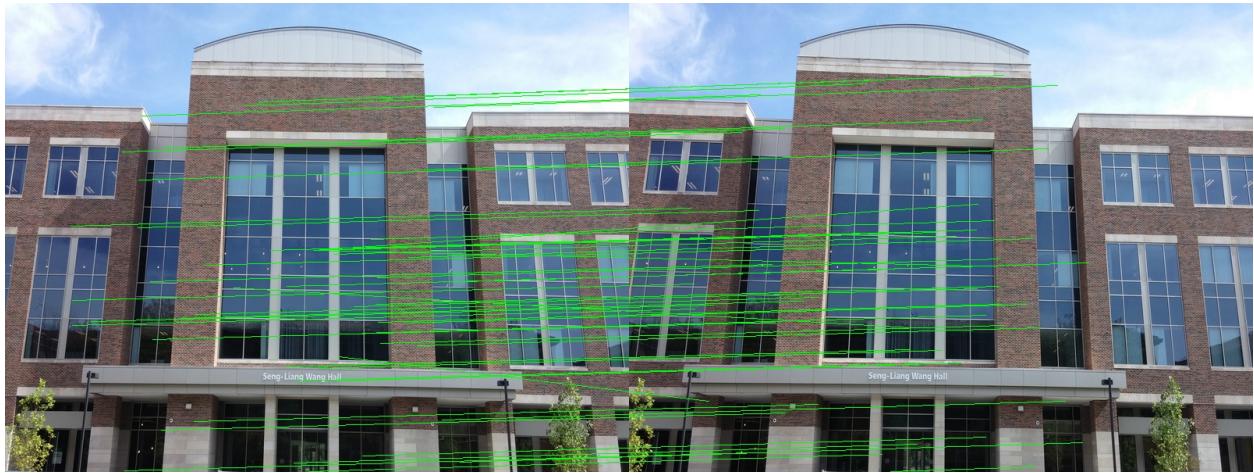
HARRIS (thresh = 0.990, sigma = 1.1414)



HARRIS (thresh = 0.993, sigma = 1.1414)



HARRIS (thresh = 0.99, sigma = 2.0)



Pair 2

Original Images

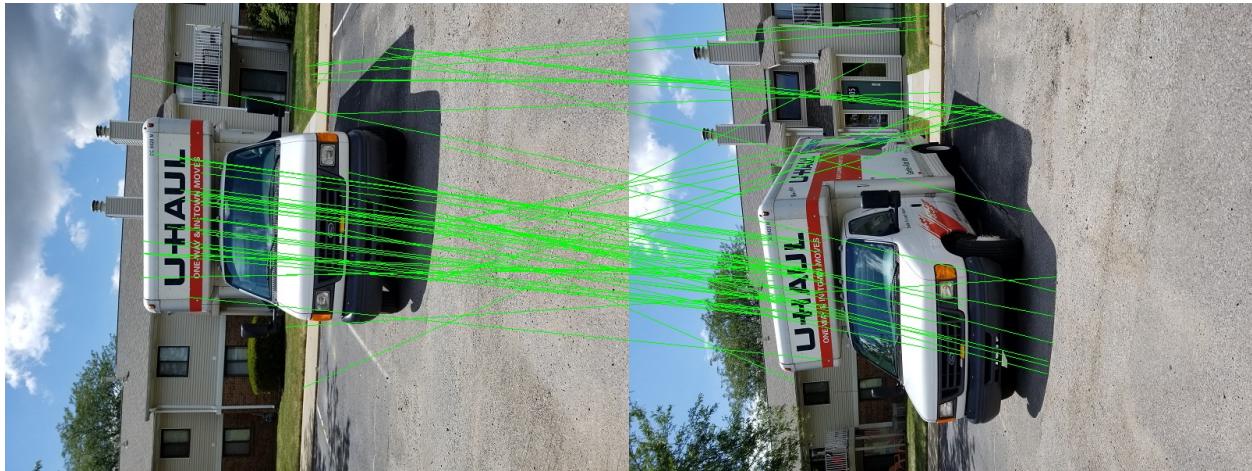


(a) Image 1



(b) Image 2

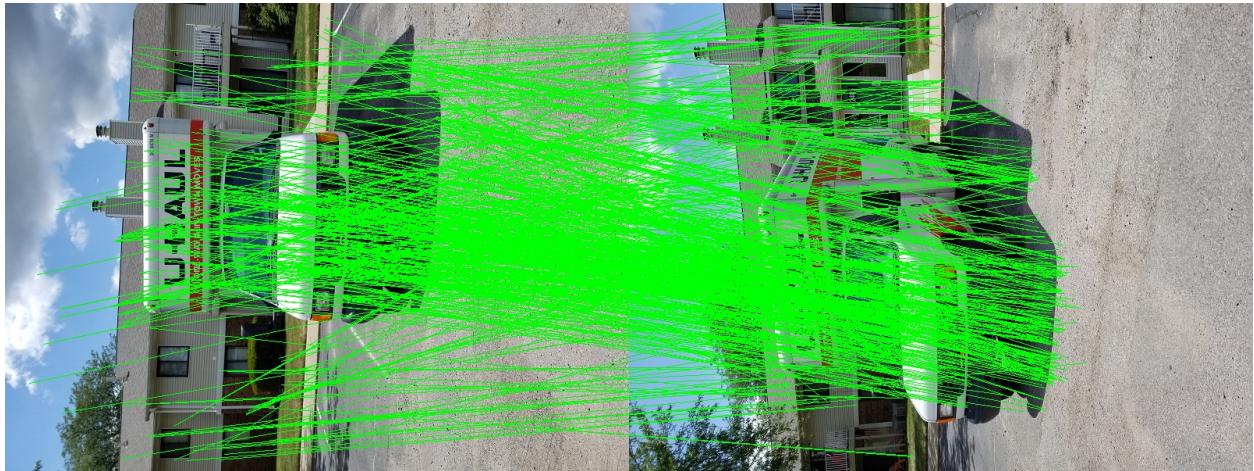
SIFT (thresh = 0.98)



SIFT (thresh = 0.99)



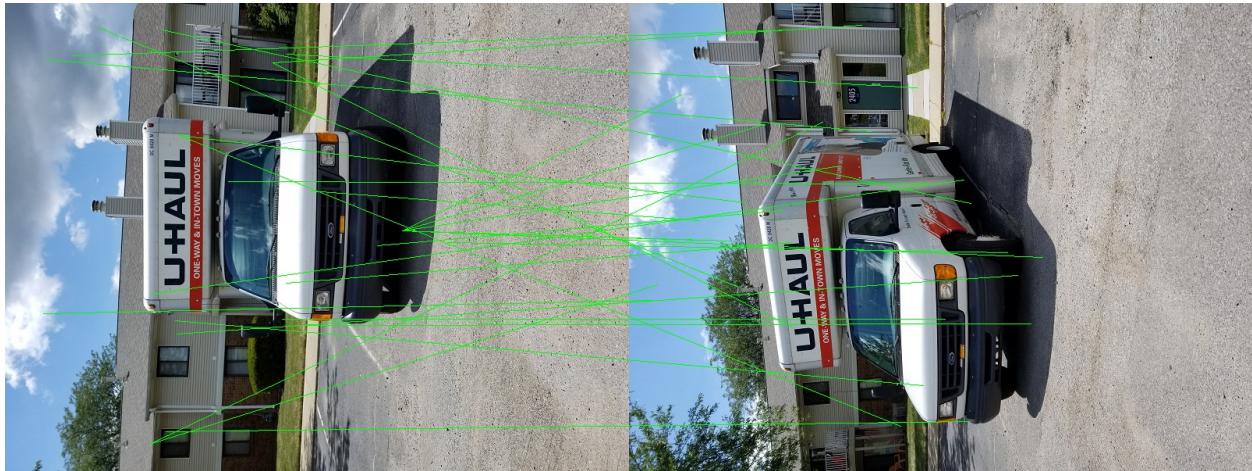
SURF (thresh = 0.99)



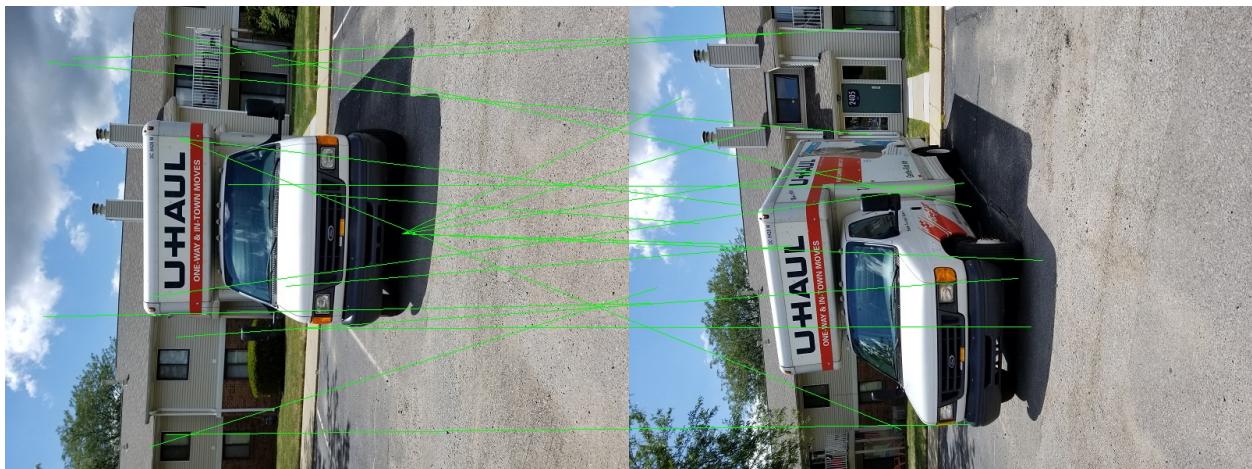
SURF (thresh = 0.9990)



HARRIS (thresh = 0.910, sigma = 1.1414)



HARRIS (thresh = 0.92, sigma = 1.1414)



HARRIS (thresh = 0.94, sigma = 2.0)



Pair 3

Original Images

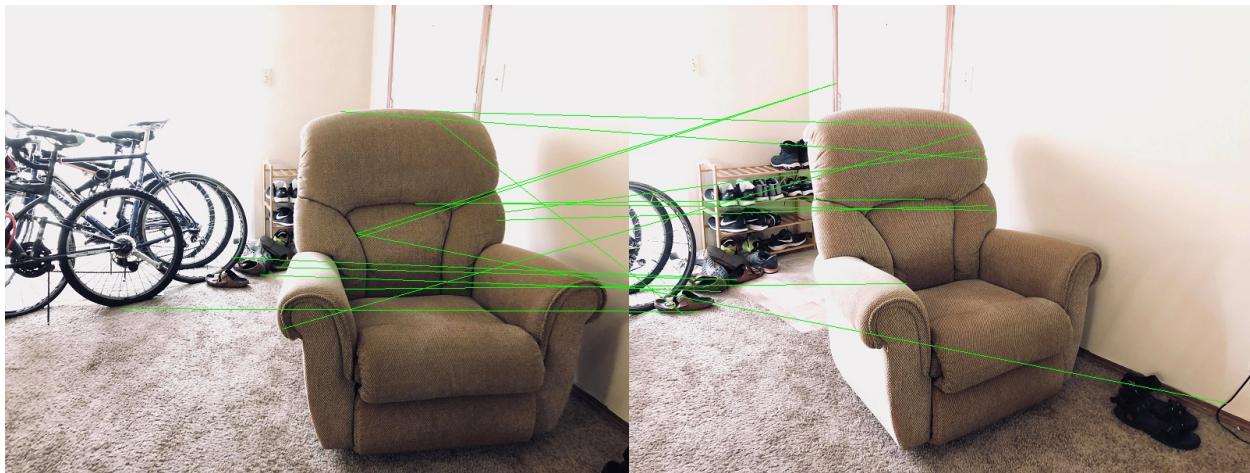


(a) Image 1



(b) Image 2

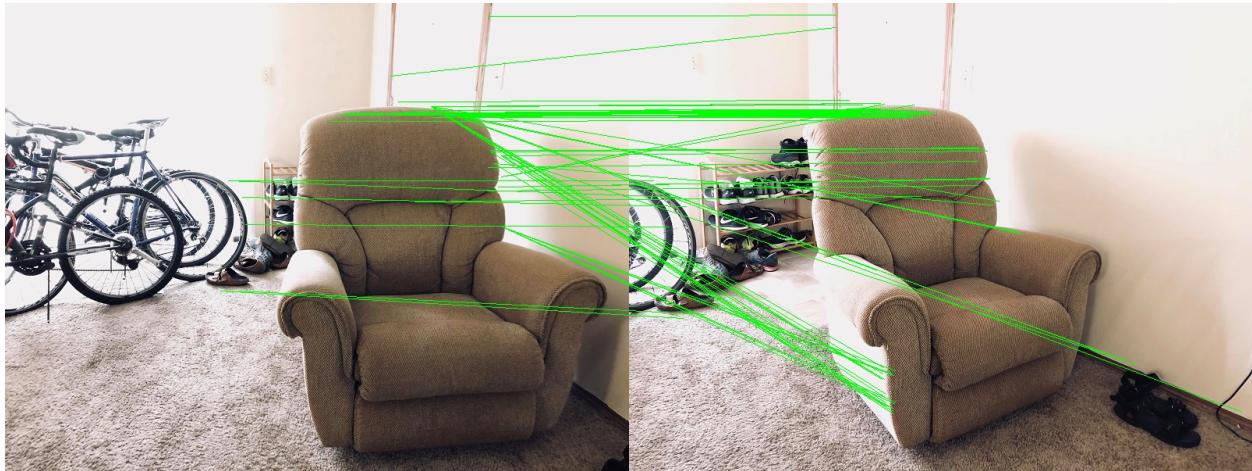
SIFT (thresh = 0.97)



SIFT (thresh = 0.99)



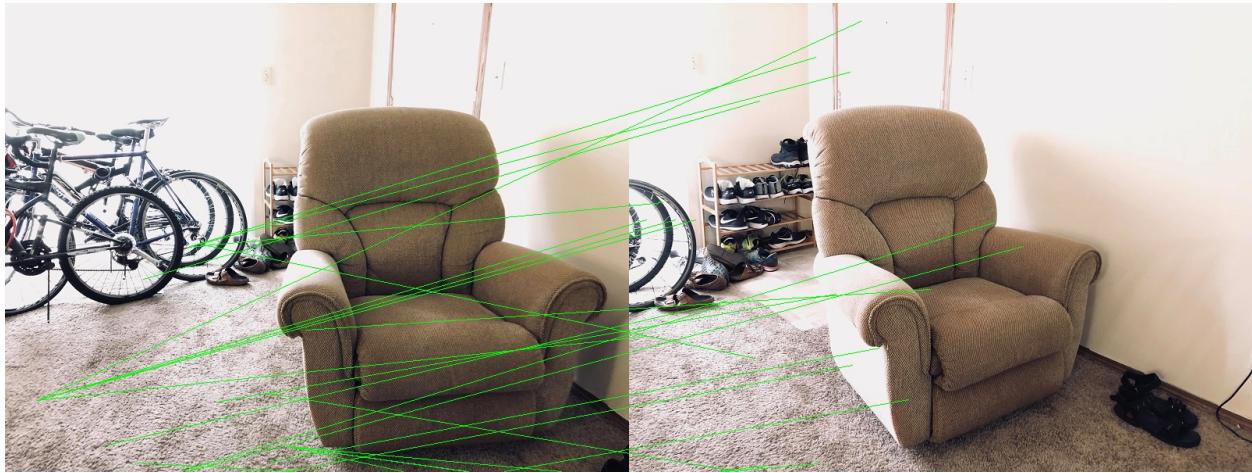
SURF (thresh = 0.996)



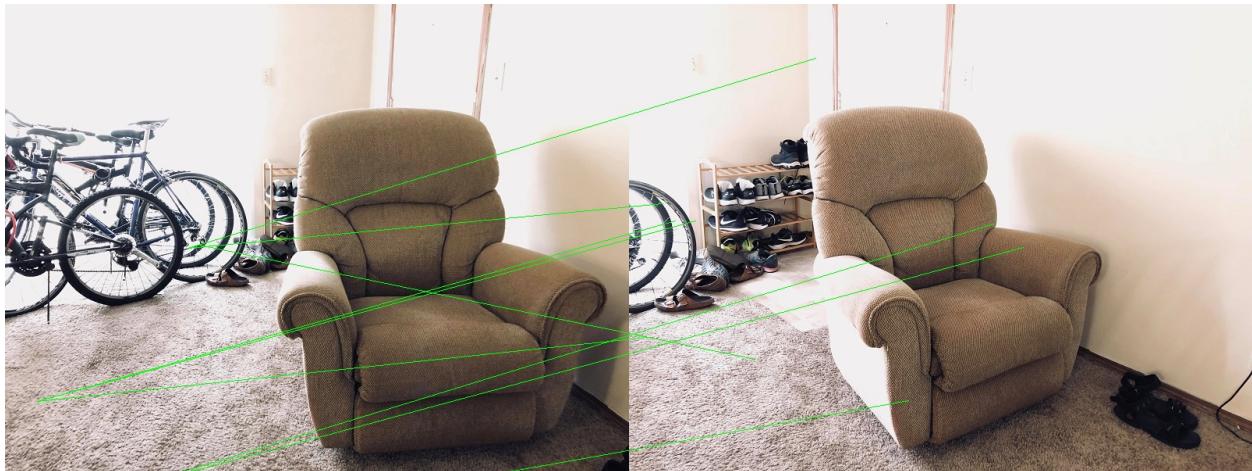
SURF (thresh = 0.9900)



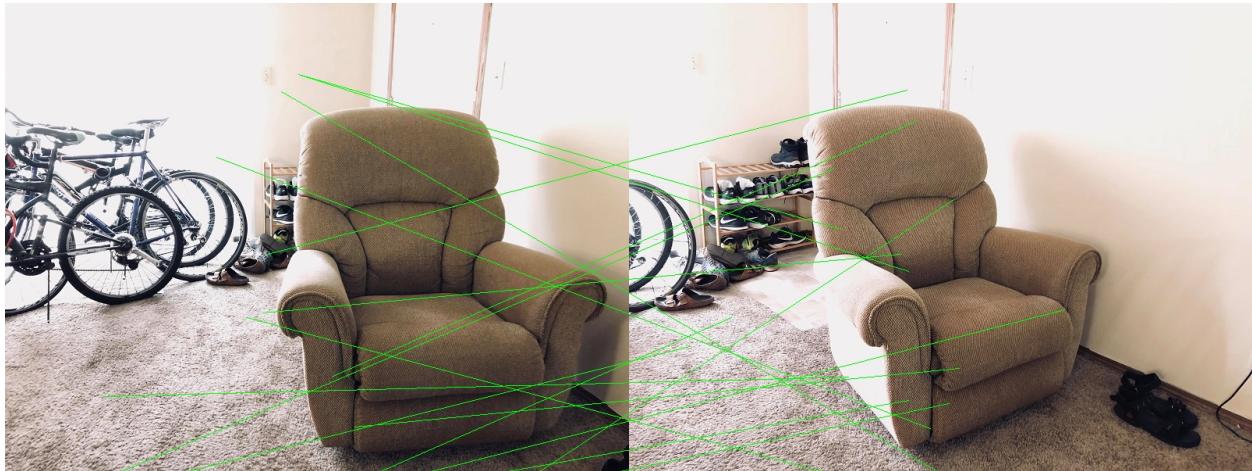
HARRIS (thresh = 0.910, sigma = 1.1414)



HARRIS (thresh = 0.94, sigma = 1.1414)



HARRIS (thresh = 0.92, sigma = 2.0)



Pair 4

Original Images

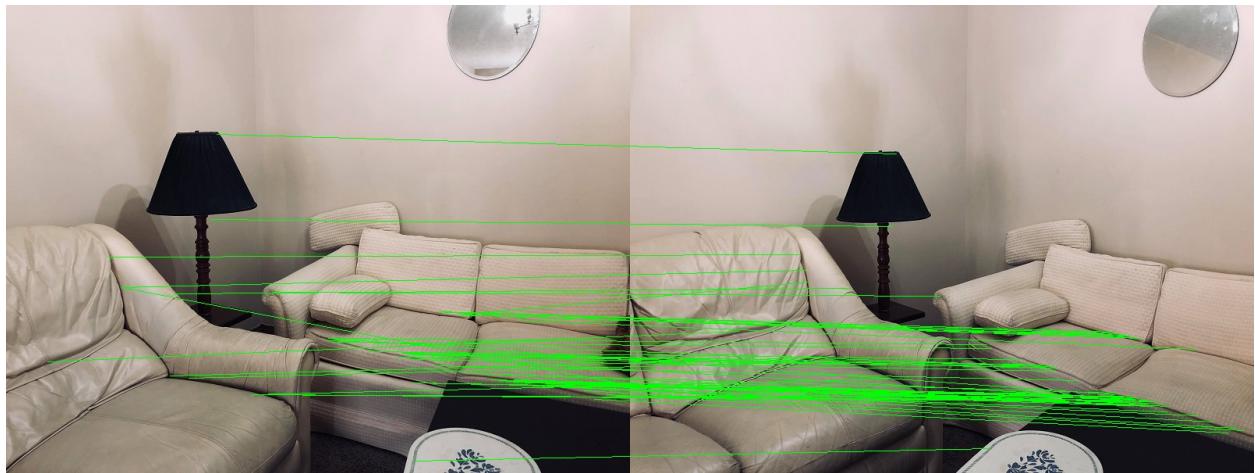


(a) Image 1

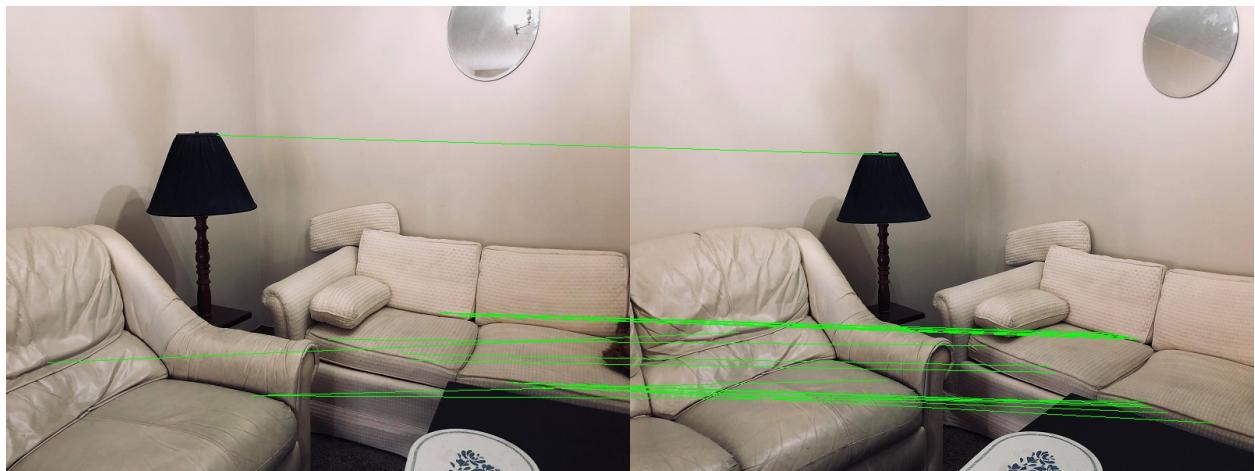


(b) Image 2

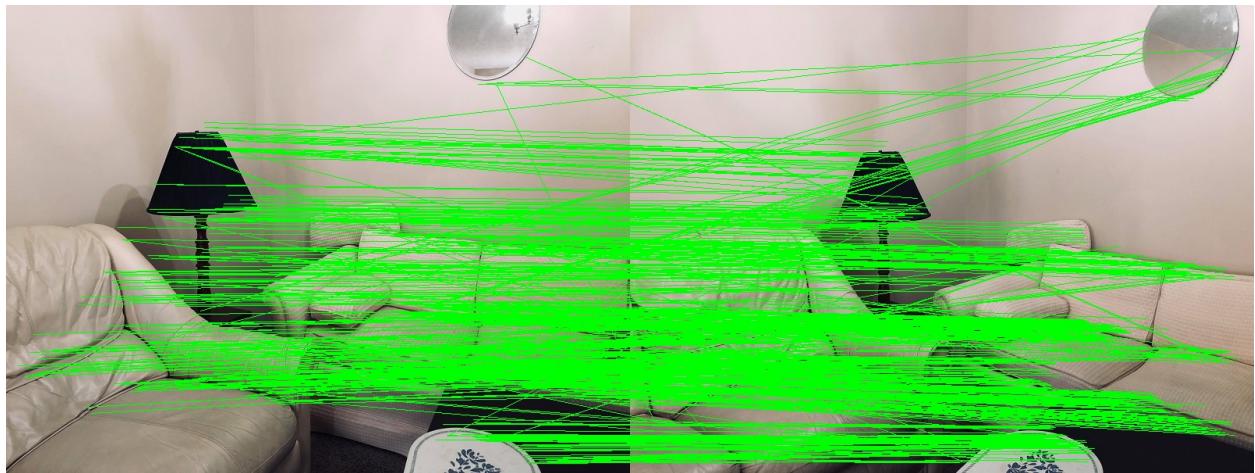
SIFT (thresh = 0.98)



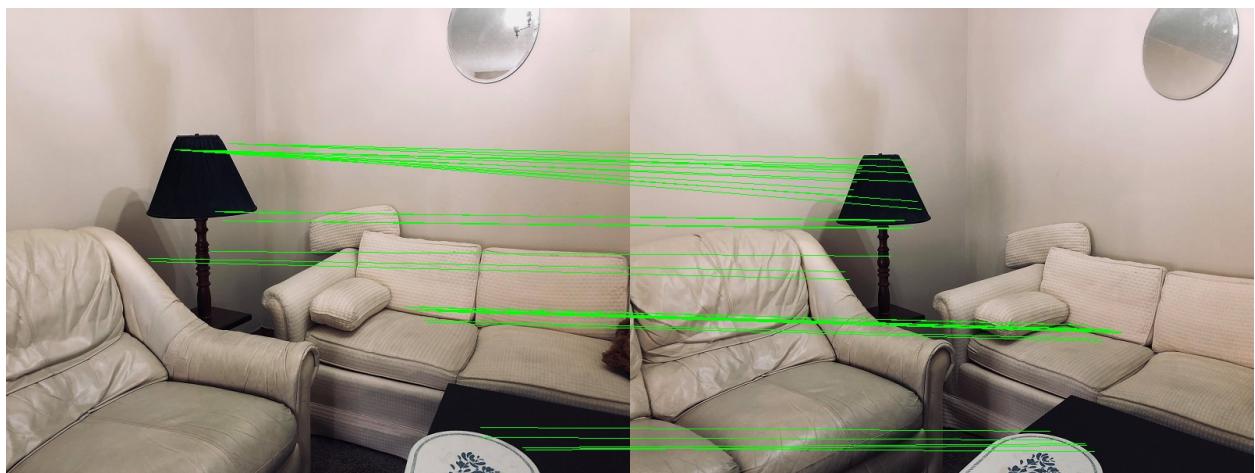
SIFT (thresh = 0.99)



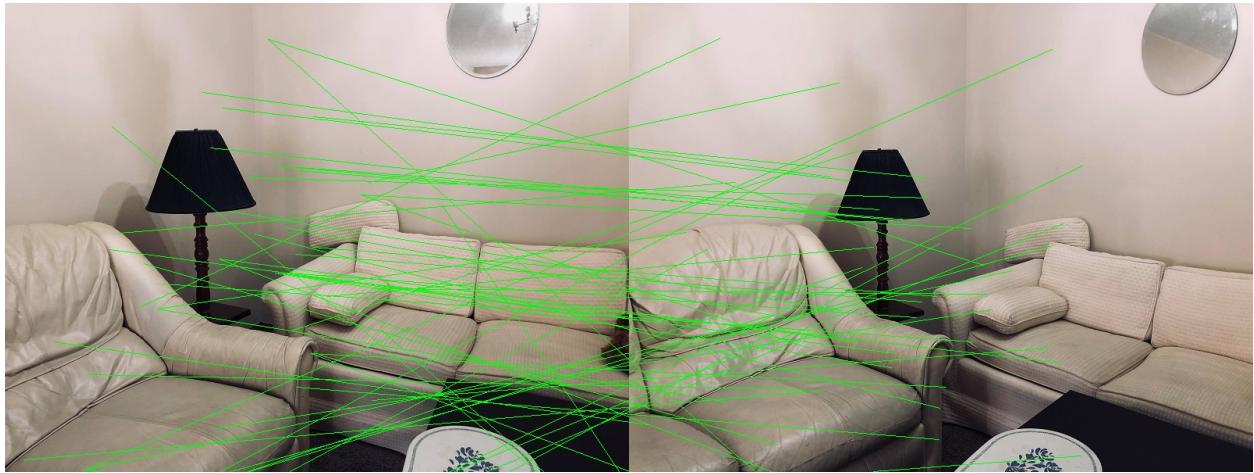
SURF (thresh = 0.99)



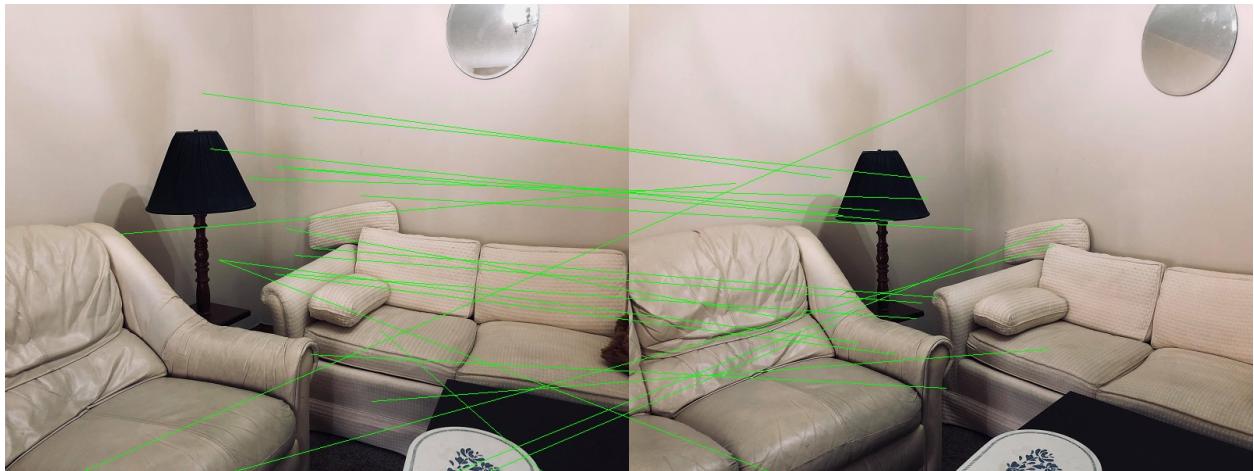
SURF (thresh = 0.9990)

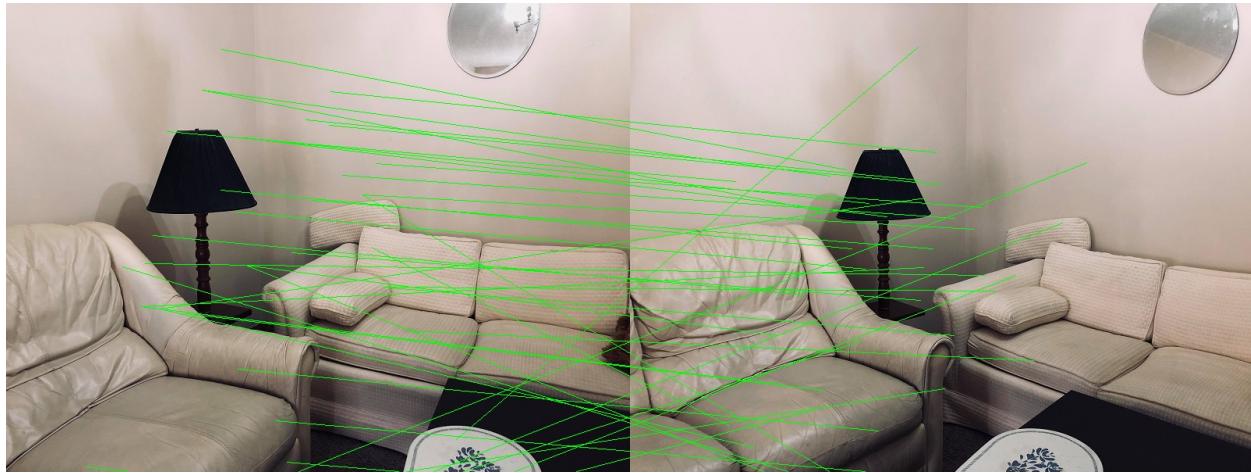


HARRIS (thresh = 0.910, sigma = 1.1414)



HARRIS (thresh = 0.94, sigma = 1.1414)



HARRIS (thresh = 0.94, sigma = 2.0)**Source Code**

```

import cv2
import numpy as np
import time
from os.path import join, basename, splitext, dirname
5
def harris(img, sigma = 1.414, thresh = 0.08):
    # img is gray scale image

    # Size of Gaussian Kernel
10   K = int(2 * np.floor(3 * 1.414 * sigma) + 1)
    # img = cv2.GaussianBlur(img, (K, K), 0)

    img = cv2.GaussianBlur(img, None, sigma)
15   dx_img = cv2.Sobel(img, cv2.CV_64F, 1, 0, ksize=K)
    dy_img = cv2.Sobel(img, cv2.CV_64F, 0, 1, ksize=K)

    img_height, img_width = img.shape[:2]
20
    Nh = int(5 * sigma)
    if (Nh % 2 == 0): Nh += 1

    df_idx = (Nh-1)/2

25
    H = np.zeros((img_height-2*df_idx, img_width-2*df_idx))

    nms_win_size = 31 # non maximum suppression window size.
    nms_half = (nms_win_size-1)/2

    desc_size = 8 # window size is 31
30
    kps = []
    features = []

```

```

35   for zr_idx, row_idx in enumerate(range(df_idx, img_height - df_idx)):
40       row_ids = range(row_idx-df_idx, row_idx+df_idx)
        for zc_idx, col_idx in enumerate(range(df_idx, img_width - df_idx)):
            col_ids = range(col_idx-df_idx, col_idx+df_idx)
            M = img[row_ids, col_ids]
            dMx = dx_img[row_ids, col_ids]
            dMy = dy_img[row_ids, col_ids]

45       c11 = np.sum(np.sum(dMx**2))
           c22 = np.sum(np.sum(dMy**2))
           c12 = np.sum(np.multiply(dMx, dMy).flatten())
           c21 = c12

           trace = c11 + c22
           det = c11 * c22 - c12 * c21

50       if(trace < 1e-5): continue

           beta = det / (trace * trace) # Min = 0, Max = 0.25. Higher the better.

55       if(beta > thresh):
           H[zr_idx, zc_idx] = beta

## Non Maximum suppression
60       for zr_idx, row_idx in enumerate(range(df_idx, img_height - df_idx)):
           for zc_idx, col_idx in enumerate(range(df_idx, img_width - df_idx)):

65           # if did not pass previous threshold, continue
           if(H[zr_idx, zc_idx] == 0): continue

           # Boundary conditions. Ignore points on the boundary.
           if(zr_idx - nms_half < 0 or zc_idx - nms_half < 0): continue
           if(zr_idx + nms_half > img_height-1 or \
              zc_idx + nms_half > img_width-1): continue

70           nms_win = H[zr_idx - nms_half:zr_idx + nms_half, \
                         zc_idx - nms_half:zc_idx + nms_half]

           # if beta is max in that window, then consider. Ignore otherwise.
           if(np.max(nms_win[:]) == H[zr_idx, zc_idx]):
               kps.append([row_idx, col_idx])
               features.append(img[row_idx-desc_size:row_idx+desc_size, \
                                 col_idx-desc_size:col_idx+desc_size].flatten().tolist())

75       return np.float32(kps), np.float32(features)

80 def extract_kps(image, ftype = 'sift', sigma = 1.414):
    # ftype - feature type 'sift' or 'surf'
    # image - np.ndarray (_ x _ x 3)
    # convert the image to grayscale
    gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
85   if(ftype.lower() == 'sift'):
```

```

descriptor = cv2.xfeatures2d.SIFT_create()
# keypoints (cv2.KeyPoint object) and features (ndarray).
(kps, features) = descriptor.detectAndCompute(image, None)
kps = np.float32([kp.pt for kp in kps])

90  elif ftype.lower() == 'surf':
    descriptor = cv2.xfeatures2d.SURF_create()
    # keypoints (cv2.KeyPoint object) and features (ndarray).
    (kps, features) = descriptor.detectAndCompute(image, None)
    kps = np.float32([kp.pt for kp in kps])

95  elif ftype.lower() == 'harris':
    start = time.time()
    (kps, features) = harris(gray, sigma = sigma)
    print 'Harris corners: %.02f secs' %(time.time()-start)

100 # kps: ndarray (_ x 2); features: ndarray (_ x 128)
     return (kps, features)

def mean_normalize(M, axis = 0):
    # Mean normalize matrix M (_ x k) in rows
    if(axis == 1): M = M.transpose() # (k x _)
    M -= np.mean(M, axis = 0)

    norms = np.linalg.norm(M, axis = 0)
    norms[norms == 0.0] = 1e-10

110 M /= norms
    if(axis == 1): M = M.transpose()
    return M

115 def dist_mat_vec(M, vec, method = 'ncc'):
    # M : ndarray (_ x k); vec: (1 x k)
    if(method.lower() == 'ssd'):
        return np.linalg.norm(M - vec, axis = 1)
    elif(method.lower() == 'ncc'):
        # Mean Normalizing rows of M
        M = mean_normalize(M, axis = 1)
        # Mean normalizing vec
        vec = vec - np.mean(vec)
        vect = vec / np.linalg.norm(vec)
        return np.dot(M, vect)
    elif(method.lower() == 'dot'):
        return np.dot(M, vec)

120 def dist_mat_mat(M1, M2, method = 'ncc'):
    # M1, M2 --> ndarray (y1 x k) and (y2 x k)
    D = np.zeros((M1.shape[0], M2.shape[0]))
    if(method.lower() == 'ncc'):

130        M1 = mean_normalize(M1, axis = 1)
        M2 = mean_normalize(M2, axis = 1)
        method = 'dot'
        for idx2 in range(M2.shape[0]):
            D[:, idx2] = dist_mat_vec(M1, M2[idx2, :], method = method)
    return D

```

```

140 def filter_kps(kpA, kpB, featuresA, featuresB, method = 'ncc', thresh = 0.97):
141     print len(featuresA), len(featuresB)
142     if(method.lower() == 'ncc'):
143         featuresA = mean_normalize(featuresA, axis = 1)
144         featuresB = mean_normalize(featuresB, axis = 1)
145         method = 'dot'
146
147     matches = []
148     for idxB in range(featuresB.shape[0]):
149         temp = dist_mat_vec(featuresA, featuresB[idxB, :], method = method)
150         temp[temp<thresh] = 0.0
151         if(np.max(temp) == 0.0): continue
152         else: matches.append((np.argmax(temp), idxB))
153
154     return matches
155
156 def draw_matches(imageA, imageB, kpsA, kpsB, matches):
157     # initialize the output visualization image
158     (hA, wA) = imageA.shape[:2]
159     (hB, wB) = imageB.shape[:2]
160     vis = np.zeros((max(hA, hB), wA + wB, 3), dtype="uint8")
161     vis[0:hA, 0:wA] = imageA
162     vis[0:hB, wA:] = imageB
163
164     # loop over the matches
165     for queryIdx, trainIdx in matches:
166         # only process the match if the keypoint was successfully
167         # matched
168         # draw the match
169         ptA = (int(kpsA[queryIdx][0]), int(kpsA[queryIdx][1]))
170         ptB = (int(kpsB[trainIdx][0]) + wA, int(kpsB[trainIdx][1]))
171         cv2.line(vis, ptA, ptB, (0, 255, 0), 1)
172
173     # return the visualization
174     return vis
175
176 def run(image1_path, image2_path, ftype = 'sift', method = 'ncc', \
177     thresh = 0.97, sigma = 1.414):
178     img1 = cv2.imread(image1_path)
179     img2 = cv2.imread(image2_path)
180
181     kps1, features1 = extract_kps(img1, ftype = ftype, sigma = sigma)
182     kps2, features2 = extract_kps(img2, ftype = ftype, sigma = sigma)
183
184     start = time.time()
185     matches = filter_kps(kps1, kps2, features1, features2, method, thresh = thresh)
186     print 'Filter Kps: %.02f secs'%(time.time()-start)
187
188     vis = draw_matches(img1, img2, kps1, kps2, matches)
189
190     fname = splitext(basename(image1_path))[0] + '_' + splitext(basename(image2_path))[0]
191     fname = fname + '_' + str(ftype) + '_' + str(int(sigma*1000)) + '_' + method + '.jpg'

```

```
195     fname_path = join(dirname(img1_path), fname)
         print 'Writing to: ', fname_path

200     cv2.imwrite(fname_path, vis)
         cv2.imshow('Visualization', vis)
         cv2.waitKey(0)

205     ##### MAIN #####
#      MAIN      #
##### MAIN #####
#      MAIN      #
##### MAIN #####
#      MAIN      #

210     if __name__ == '__main__':
         img1_path = 'pair3/1.jpg'
         img2_path = 'pair3/2.jpg'

         ftype = 'sift'
         method = 'ncc'
         thresh = 0.97 # SURF 0.9999

         vis = run(img1_path, img2_path, ftype = ftype, method=method, thresh = thresh)
```